

Complete OpenTelemetry Observability Project Guide

Executive Summary

This comprehensive guide walks through deploying the OpenTelemetry (OTel) Astronomy Shop demonstration application—a production-grade microservice architecture for learning observability. This project showcases how to collect telemetry data (metrics, logs, and traces) from complex distributed systems using OpenTelemetry, then visualize and analyze them using industry-standard tools like Prometheus, Grafana, Jaeger, and OpenSearch[1].

The guide covers architecture understanding, step-by-step deployment on AWS, configuration management, data analysis, and integration with multiple observability backends. By completing this project, you'll demonstrate enterprise-level DevOps and observability expertise[2].

Table of Contents

1. [Understanding OpenTelemetry](#)
 2. [Architecture Overview](#)
 3. [Prerequisites and Setup](#)
 4. [AWS Environment Configuration](#)
 5. [Application Deployment](#)
 6. [Monitoring and Observability](#)
 7. [Feature Flags and Error Injection](#)
 8. [Advanced Configuration](#)
 9. [Troubleshooting Guide](#)
 10. [Best Practices](#)
-

Understanding OpenTelemetry

What is OpenTelemetry?

OpenTelemetry (OTel) is an open standard for data collection that is vendor-agnostic and maintained by the Cloud Native Computing Foundation (CNCF). As the second-most active CNCF project after Kubernetes, it provides a unified approach to observability[1].

Key Characteristics:

- **Vendor-agnostic:** Works with any monitoring platform (Prometheus, Grafana, Datadog, New Relic, Splunk, Dynatrace, etc.)
- **Language-agnostic:** Supports 10+ programming languages with native SDKs
- **Standardized:** Defines semantic conventions for consistent telemetry data
- **Open-source:** Community-driven development and contributions

The Three Pillars of Observability

Observability relies on three complementary data types to provide complete system visibility[3]:

1. Metrics

Purpose: Track system performance and health over time

Characteristics:

- Numerical measurements aggregated at specific intervals
- Include CPU usage, memory consumption, request rate, latency, error rate
- Stored efficiently due to time-series nature
- Enable trend analysis and anomaly detection

Example:

```
service_cpu_usage: 65%
http_request_duration: 150ms
error_rate: 2.5%
```

2. Logs

Purpose: Capture detailed, moment-specific information about system events

Characteristics:

- Text-based or structured records of events and errors
- Include timestamps, severity levels, and contextual information
- Provide detailed debugging information
- Historical record of system behavior

Example:

```
{
  "timestamp": "2024-12-04T10:08:00Z",
  "level": "ERROR",
  "service": "payment-service",
  "message": "Payment processing failed",
  "error": "Connection timeout",
  "trace_id": "4bf92f3577b34da6a3ce929d0e0e4736"
}
```

3. Traces

Purpose: Show the complete journey of requests across microservices

Characteristics:

- Map request flow through distributed systems
- Identify latency sources and bottlenecks
- Show service dependencies and communication patterns
- Enable root cause analysis

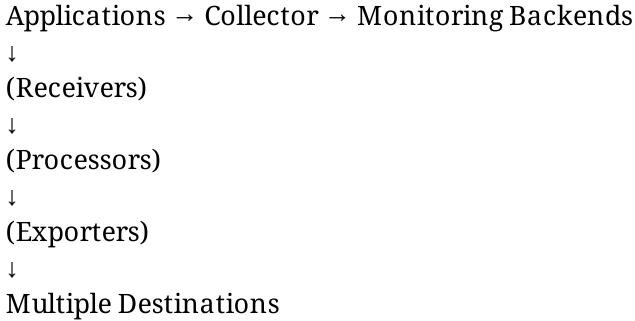
Example:

Request Flow: LoadGenerator → FrontEnd Proxy → Frontend → Ad Service

Span 1: LoadGenerator (0-50ms)
Span 2: FrontEnd Proxy (50-100ms)
Span 3: Frontend (100-150ms)
Span 4: Ad Service (150-250ms, ERROR)

OpenTelemetry Collector Architecture

The OTel Collector is a standalone service that receives, processes, and exports telemetry data[4]:



Key Components

Receivers (Data Ingestion)

- OTLP: OpenTelemetry Protocol (native, recommended)
- Jaeger: Distributed tracing format
- Prometheus: Metrics scraping
- OpenCensus: Legacy tracing format
- Many others for specific use cases

Processors (Data Enhancement)

- Batch: Group data for efficiency
- Attributes: Add metadata to spans
- Resource Detection: Automatically add resource information
- Sampling: Reduce data volume intelligently
- Memory Limiter: Prevent resource exhaustion

Exporters (Data Delivery)

- OTLP: To other OTel services or backends
- Prometheus: For metrics
- Jaeger: For traces
- OpenSearch: For logs
- Splunk, Datadog, New Relic, etc.

Extensions (Optional Features)

- Health Check: Collector status monitoring
- Authentication: Secure communication
- zPages: In-process diagnostics

Architecture Overview

OpenTelemetry Astronomy Shop Demo

The demo application is a cloud-native e-commerce system built to showcase OpenTelemetry capabilities across multiple programming languages and services[1].

Microservices Overview

The application consists of 18+ microservices built in different languages:

Service	Language	Purpose
Frontend	JavaScript/TypeScript	Web interface for the store
Frontend Proxy	Go	API gateway and request routing
Ad Service	Java	Advertisement recommendations
Cart Service	.NET	Shopping cart management
Checkout Service	Go	Order processing
Payment Service	JavaScript/Node.js	Payment processing
Shipping Service	Go	Shipping logistics
Product Catalog	Go	Product information database
Currency Service	Node.js	Currency conversion
Recommendation Service	Python	Product recommendations
Accounting Service	.NET	Financial tracking
Email Service	Python	Email notifications
Order Service	Java	Order management
Flag Service	Go	Feature flag management
Kafka	Message Queue	Event streaming

Technology Stack

Telemetry Collection:

- OpenTelemetry Collector
- OpenTelemetry SDKs (one per language)

Observability Backends (Default):

- Prometheus: Metrics storage and querying
- Grafana: Metrics visualization and dashboards
- Jaeger: Distributed tracing
- OpenSearch: Log aggregation and search

Additional Components:

- Locust: Load generation for realistic traffic
- Docker Compose: Container orchestration
- Kafka: Event streaming

Data Flow Architecture

Microservices

↓

OpenTelemetry Instrumentation

↓

OpenTelemetry Collector

↓

|→ Prometheus (Metrics)
|→ Jaeger (Traces)
|→ OpenSearch (Logs)
└→ Other Backends (Optional)

↓

Visualization & Alerting

|→ Grafana Dashboards
|→ Jaeger UI
|→ OpenSearch Dashboards
└→ Alert Rules

Prerequisites and Setup

System Requirements

Minimum Requirements:

- AWS EC2 instance with **t2.xlarge** (16GB RAM, 4 vCPU)
 - Kafka requires significant memory
 - Total data collection requires processing power
 - Free tier (t2.micro with 1GB RAM) is insufficient

Estimated Costs:

- t2.xlarge: ~\$0.15/hour

- Storage: 15GB EBS at ~\$1.50/month
- Data transfer: Minimal for local testing
- **Total:** ~\$5-10 for 24-hour project

Required Tools

On Local Machine:

- SSH client (built-in on Linux/macOS, PuTTY on Windows)
- AWS account with EC2 permissions
- Terminal/Command prompt

On AWS EC2 Instance:

- Docker (container runtime)
- Docker Compose (container orchestration)
- Git (repository cloning)
- Ubuntu 22.04 LTS operating system

Knowledge Prerequisites

- Basic Linux command-line operations
- Understanding of microservices architecture
- Familiarity with containers and Docker
- Basic networking concepts (ports, protocols)
- No need for deep knowledge of individual languages

GitHub Repository

Official Repository:

<https://github.com/open-telemetry/opentelemetry-demo>

Documentation:

<https://opentelemetry.io/docs/demo/>

AWS Environment Configuration

Step 1: Launch EC2 Instance

Access AWS Console

1. Log in to AWS Management Console
2. Navigate to EC2 Dashboard
3. Click **Instances**
4. Click **Launch Instance**

Instance Configuration

Name: otel-observability-project

Operating System:

- Ubuntu 22.04 LTS (Amazon Machine Image)
- Free tier eligible, widely supported

Instance Type: t2.xlarge

- 16 GB RAM
- 4 vCPU
- Not eligible for free tier (charges apply)

Key Pair:

- Create new or use existing SSH key
- Save securely (cannot be recovered)
- Required for SSH access

Security Group Rules:

Inbound Rules:

- SSH (Port 22): From 0.0.0.0/0 (your IP recommended)
- HTTP (Port 80): From 0.0.0.0/0 (application)
- HTTP (Port 8080): From 0.0.0.0/0 (load generator; Jaeger, Grafana)
- Custom TCP (Port 4317): From 0.0.0.0/0 (OTel Collector GRPC)
- Custom TCP (Port 4318): From 0.0.0.0/0 (OTel Collector HTTP)

Outbound Rules:

- All traffic allowed (default)

Storage:

- EBS Volume: 15 GB gp3
- Sufficient for Docker images and logs

Step-by-Step Launch

Instance launches (takes 30-60 seconds)

Wait for status to change from "Pending" to "Running"

Wait for Status Checks to pass (2/2 checks)

Step 2: Connect to Instance

Option A: Using EC2 Instance Connect (Easiest)

1. Select instance in AWS Console
2. Click **Connect** button
3. Choose **EC2 Instance Connect** tab
4. Click **Connect** (opens browser terminal)

Option B: Using SSH from Terminal

After saving key pair

```
chmod 400 ~/Downloads/your-key-pair.pem
```

Connect via SSH

```
ssh -i ~/Downloads/your-key-pair.pem ubuntu@<PUBLIC_IP>
```

Example:

```
ssh -i ~/Downloads/your-key-pair.pem ubuntu@54.123.456.789
```

Step 3: Update System Packages

Run on EC2 instance

```
sudo apt update -y  
sudo apt upgrade -y
```

This ensures all packages are current

Prevents version conflicts and security issues

Step 4: Install Docker and Docker Compose

Install Docker

```
sudo apt install -y dockerio docker-compose
```

Verify installation

```
docker --version  
docker-compose --version
```

Add current user to docker group (optional, for convenience)

```
sudo usermod -aG docker ubuntu
```

Note: Docker Compose comes bundled with modern Docker installations as docker compose (newer) or separate docker-compose (older).

Application Deployment

Step 1: Clone OpenTelemetry Demo Repository

Navigate to home directory

```
cd ~
```

Clone the official repository

```
git clone https://github.com/open-telemetry/opentelemetry-demo.git
```

Navigate to project

```
cd opentelemetry-demo
```

List contents

```
ls -la
```

Key files:

- docker-compose.yml: Main orchestration file
- src/: Source code for all microservices

- docs/: Documentation and guides

Step 2: Understanding Docker Compose File

The docker-compose.yml file defines all services, networks, and configurations:

```
version: '3'
```

Global logging configuration

```
x-default-logging: &default-logging
  driver: "json-file"
  options:
    max-size: "5m"
    max-file: "2"
```

Network definition for service communication

```
networks:
  default:
    name: opentelemetry-demo
    driver: bridge
```

Service definitions (18+ services)

```
services:
```

Example: Accounting Service (.NET)

```
accountingservice:
  image: ghcr.io/open-telemetry/demo:latest-accountingservice
  container_name: accountingservice
  depends_on:
    otelcol:
      condition: service_started
  environment:
    - OTEL_EXPORTER_OTLP_ENDPOINT=http://otelcol:4317
    - OTEL_RESOURCE_ATTRIBUTES=service.name=accountingservice
  ports:
    - "8200:8200"
  logging: *default-logging
```

Example: Frontend (JavaScript/Node.js)

```
frontend:  
image: ghcr.io/open-telemetry/demo:latest-frontend  
container_name: frontend  
ports:  
- "3000:3000"  
environment:  
- PUBLIC_OTEL_EXPORTER_OTLP_TRACES_ENDPOINT=http://localhost:4318/v1/traces  
depends_on:  
- frontendproxy  
logging: *default-logging
```

OpenTelemetry Collector

```
otelcol:  
image: otel/opentelemetry-collector-k8s:latest  
container_name: otelcol  
command: [ "--config=/etc/otel-collector-config.yml" ]  
volumes:  
- ./src/otelcollector/otel-collector-config.yml:/etc/otel-collector-config.yml  
ports:  
- "4317:4317" # OTLP GRPC receiver  
- "4318:4318" # OTLP HTTP receiver  
- "9411:9411" # Zipkin receiver  
environment:  
- GOGC=80  
depends_on:  
- jaeger  
- prometheus  
logging: *default-logging
```

Prometheus for metrics

```
prometheus:  
image: prom/prometheus:latest  
container_name: prometheus  
command:  
- "--config.file=/etc/prometheus/prometheus.yml"  
- "--storage.tsdb.path=/prometheus"  
volumes:  
- ./src/prometheus/prometheus.yml:/etc/prometheus/prometheus.yml  
ports:  
- "9090:9090"  
logging: *default-logging
```

Grafana for visualization

```
grafana:  
image: grafana/grafana:latest  
container_name: grafana  
environment:  
- GF_AUTH_ANONYMOUS_ENABLED=true  
- GF_AUTH_ANONYMOUS_ORG_ROLE=Admin  
volumes:  
- ./src/grafana/provisioning:/etc/grafana/provisioning  
- ./src/grafana/dashboards:/var/lib/grafana/dashboards  
ports:  
- "3001:3000"  
depends_on:  
- prometheus  
logging: *default-logging
```

Jaeger for tracing

```
jaeger:  
image: jaegertracing/all-in-one:latest  
container_name: jaeger  
ports:  
- "16686:16686" # Jaeger UI  
- "4317:4317" # OTLP GRPC receiver  
- "9411:9411" # Zipkin receiver  
logging: *default-logging
```

Kafka for event streaming

```
kafka:  
image: confluentinc/cp-kafka:7.5.0  
container_name: kafka  
depends_on:  
- zookeeper  
environment:  
KAFKA_BROKER_ID: 1  
KAFKA_ZOOKEEPER_CONNECT: zookeeper:2181  
KAFKA_ADVERTISED_LISTENERS: PLAINTEXT://kafka:9092  
KAFKA_OFFSETS_TOPIC_REPLICATION_FACTOR: 1  
logging: *default-logging
```

Zookeeper for Kafka coordination

```
zookeeper:  
image: confluentinc/cp-zookeeper:7.5.0  
container_name: zookeeper  
environment:  
ZOOKEEPER_CLIENT_PORT: 2181  
logging: *default-logging
```

... (additional services defined similarly)

Key Concepts:

Section	Purpose
version	Docker Compose file format
x-default-logging	Logging template reused by services
networks	Define service communication
services	Container definitions
image	Docker image to use
container_name	Hostname within network
ports	Port mappings: HOST:CONTAINER
environment	Environment variables
depends_on	Service startup order
volumes	Mount configurations
logging	Log collection settings

Step 3: Deploy Application

Ensure you're in the project directory

```
cd ~/opentelemetry-demo
```

Start all services

```
sudo docker-compose up -d
```

The **-d** flag runs in background

Docker Compose will:

1. Pull all images from registry
2. Create network: opentelemetry-demo
3. Start all 18+ services in order
4. Wait for health checks

This takes 2-5 minutes on first run

Step 4: Verify Deployment

Check running containers

```
sudo docker ps
```

Expected output: 18-20 containers running

NAME STATUS PORTS

accountingservice Up 2 minutes 8200/tcp

adservice Up 2 minutes 9555/tcp

cartservice Up 2 minutes 7070/tcp

...

Check logs for specific service

sudo docker logs <service_name>

Example: Check Collector logs

sudo docker logs otelcol

Check for errors

sudo docker logs otelcol | grep -i error

View all logs

sudo docker-compose logs -f

Exit with Ctrl+C

Step 5: Access Application Endpoints

Once deployment succeeds, access applications via instance public IP:

Core Endpoints:

Service	URL	Purpose
Frontend	http://<IP>:80	E-commerce store frontend
Grafana	http://<IP>:3001/grafana	Metrics dashboards
Jaeger	http://<IP>:16686	Distributed tracing UI
Prometheus	http://<IP>:9090	Metrics database
Load Generator	http://<IP>:8089/	Generate synthetic traffic
Kafka	kafka:9092	Event broker (internal)
OpenSearch	http://<IP>:9200	Log storage (internal)

Finding Public IP:

Option 1: From AWS Console

Click instance → Note IPv4 Public IP

Option 2: From terminal

curl <http://169.254.169.254/latest/meta-data/public-ipv4>

Test Connectivity:

From your local machine

curl http://<PUBLIC_IP>:80

Should return HTML from frontend

If connection refused, wait 30 seconds and retry

Monitoring and Observability

Accessing Grafana Dashboards

Login and Navigation

1. Open browser: `http://<PUBLIC_IP>:3001`
2. Default credentials (no login required if anonymous enabled):
 - Username: admin
 - Password: admin

Available Dashboards

1. Demo Dashboard (Main)

- Overview of all microservices
- Request rates and latency
- Error rates per service
- Resource utilization

2. Span Metrics Demo Dashboard

- Detailed span metrics
- Service-to-service latency
- Error analysis by operation
- Top 10 services by metrics

3. OpenTelemetry Collector Dashboard

- Collector health status
- Data processing rates
- Receiver/processor/exporter metrics
- Memory and CPU usage

4. OpenSearch Logs Dashboard

- Log aggregation and search
- Error tracking
- Performance warnings
- Structured log viewing

Example: Viewing Service Metrics

1. Click on "Demo Dashboard"
2. Select service from dropdown (e.g., "adservice")
3. View metrics:
 - Latency: P50, P95, P99
 - Error Rate: Percentage of failed requests

- Request Rate: Requests per second
 - Resource Usage: CPU, Memory
4. Time range selector: Last 5 min, 15 min, 1 hour, etc.

Understanding Jaeger Traces

Jaeger UI Navigation

1. Open browser: http://<PUBLIC_IP>:16686
2. Jaeger UI for distributed tracing

Trace Analysis

Search for Traces:

1. Service: Select microservice (e.g., "frontend", "adservice")
2. Operation: Select specific operation (get, post, etc.)
3. Tags: Filter by specific attributes (e.g., error=true)
4. Lookback: Time range to search (5 min, 1 hour, etc.)
5. Limit Results: Maximum traces to return
6. Click "Find Traces"

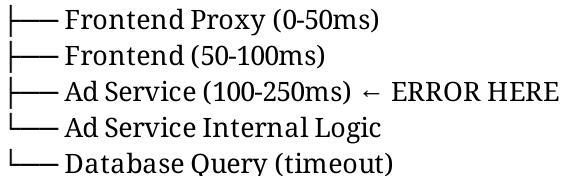
Interpreting Trace Results:

Request ID: 4bf92f3577b34da6a3ce929d0e0e4736

Duration: 250ms

Status: ERROR (indicated by red dot)

Span Breakdown:



Debugging with Traces:

1. **Identify Error Location:** Red spans indicate failures
2. **Check Error Message:** Click span → view logs and error details
3. **Analyze Dependencies:** See service call chain
4. **Measure Latency:** Find bottleneck services
5. **Compare Traces:** Use trace comparison tool for A/B analysis

Trace Data Elements

Element	Description
Trace ID	Unique identifier for complete request
Span ID	Unique identifier for service call
Parent Span ID	Links child spans to parent
Duration	Time taken for operation
Tags	Key-value metadata
Logs	Timestamped events within span
Status	Success, error, or unset

Prometheus Metrics Database

Accessing Prometheus

1. Open browser: http://<PUBLIC_IP>:9090
2. Query interface for time-series metrics

Example Queries

Request Rate (requests per second):

```
rate(http_server_request_duration_seconds_count[5m])
```

Error Rate:

```
rate(http_server_request_duration_seconds_count{status="5xx"}[5m])
/
rate(http_server_request_duration_seconds_count[5m])
```

P99 Latency:

```
histogram_quantile(0.99, rate(http_server_request_duration_seconds_bucket[5m]))
```

Service CPU Usage:

```
container_cpu_usage_seconds_total{service="adservice"}
```

Metrics Export

Download metrics in various formats:

- Prometheus text format
- JSON
- CSV

Feature Flags and Error Injection

Purpose of Feature Flags

Feature flags enable controlled error injection and scenario testing without code changes:

- Test observability tool accuracy
- Simulate production issues in controlled manner
- Validate alerting rules
- Train team on incident response

Flag Configuration File

Location:

./src/flagD/demo.flagd.json

File Structure:

```
{
  "$schema": "https://flagd.dev/schema/v0/flags.json",
  "flags": {
    "productCatalogFailure": {
      "state": "ENABLED",
      "description": "Fail product catalog service on specific product",
      "variants": {
        "on": true,
        "off": false
      },
      "defaultVariant": "off"
    },
    "recommendationServiceFailure": {
      "state": "ENABLED",
      "description": "Make recommendation service fail",
      "variants": {
        "on": true,
        "off": false
      },
      "defaultVariant": "off"
    },
    "cartServiceFailure": {
      "state": "ENABLED",
      "description": "Make cart service fail",
      "variants": {
        "on": true,
        "off": false
      },
      "defaultVariant": "off"
    },
    "adServiceHighCPU": {
      "state": "ENABLED",
      "description": "Increase CPU usage for ad service",
      "variants": {

```

```
"on": true,  
"off": false  
},  
"defaultVariant": "off"  
}  
}  
}
```

Enabling Feature Flags

Edit Configuration File

SSH into instance

```
ssh -i your-key.pem ubuntu@<PUBLIC_IP>
```

Navigate to demo directory

```
cd ~/opentelemetry-demo
```

Edit flagD configuration

```
sudo nano src/flagD/demo.flagd.json
```

Make Changes

Change "defaultVariant" from "off" to "on" for desired flags:

```
{  
  "adServiceHighCPU": {  
    "defaultVariant": "on" // Changed from "off"  
  },  
  "cartServiceFailure": {  
    "defaultVariant": "on" // Changed from "off"  
  }  
}
```

Save file: Ctrl+S, then Ctrl+X

Restart Services

Stop all services

```
sudo docker-compose down
```

Start services (with new flag configuration)

```
sudo docker-compose up -d
```

Wait for all services to be ready (~2 minutes)

```
sudo docker ps | grep -c "Up"
```

Should show 18-20 containers running

Monitoring Error Injection

Generate Traffic

1. Open Load Generator: http://<PUBLIC_IP>:8089
2. Locust interface appears
3. Start load test:
 - o Number of users: 10
 - o Spawn rate: 2 users/second
 - o Target host: <http://frontend:80>
 - o Click "Start swarming"

View Errors in Grafana

1. Open Grafana: http://<PUBLIC_IP>:3001
2. Select "Demo Dashboard"
3. Change time range to "Last 5 minutes"
4. Select service with enabled flags (e.g., "adservice" for High CPU)
5. Observe:
 - o Error rate increases
 - o Latency increases
 - o CPU usage increases
 - o Log entries appear

Find Errors in Jaeger

1. Open Jaeger: http://<PUBLIC_IP>:16686
2. Service: "adservice" (if CPU flag enabled)
3. Tags: "error=true" (filter to errors only)
4. Click "Find Traces"
5. Examine error traces:
 - o Error message
 - o Stack trace
 - o Span timing
 - o Related services

View Logs in OpenSearch

1. Open OpenSearch: `http://<PUBLIC_IP>:9200/_dashboards`
 2. Create index pattern: "logs-*"
 3. Discover tab
 4. Filter: "`service.name` = adservice"
 5. View error logs with context
-

Advanced Configuration

Switching Monitoring Backends

OpenTelemetry's vendor-agnostic nature allows using different backends:

Supported Backends

Backend	Traces	Metrics	Logs	Protocol
Jaeger	✓	-	-	OTLP
Prometheus	-	✓	-	OTLP
Datadog	✓	✓	✓	OTLP API
New Relic	✓	✓	✓	OTLP API
Splunk	✓	✓	✓	OTLP
Dynatrace	✓	✓	✓	OTLP
Elastic Stack	✓	✓	✓	OTLP

Configuration File

Location: `./src/otelcollector/otel-collector-config.yml`

This YAML file defines receivers, processors, and exporters.

Default Configuration (Jaeger/Prometheus):

```
receivers:  
otlp:  
protocols:  
grpc:  
endpoint: 0.0.0.0:4317  
http:  
endpoint: 0.0.0.0:4318
```

```
processors:  
batch:  
send_batch_size: 1024  
timeout: 5s
```

```
memory_limiter:  
check_interval: 1s  
limit_mib: 512  
spike_limit_mib: 128  
  
exporters:
```

Jaeger for traces

```
jaeger:  
endpoint: jaeger:14250  
tls:  
insecure: true
```

Prometheus for metrics

```
prometheus:  
endpoint: "0.0.0.0:8889"
```

Logging for debugging

```
logging:  
loglevel: debug  
  
extensions:  
health_check:  
endpoint: ":13133"  
  
service:  
extensions: [health_check]  
pipelines:  
traces:  
receivers: [otlp]  
processors: [batch, memory_limiter]  
exporters: [jaeger, logging]
```

```
metrics:  
receivers: [otlp, prometheus]  
processors: [batch, memory_limiter]  
exporters: [prometheus, logging]
```

Switching to Datadog (Example)

Prerequisites:

- Datadog account with API key
- Collector needs network access to Datadog endpoints

Configuration Changes:

```
exporters:  
datadog:  
api:  
key: "${DD_API_KEY}" # Set as environment variable  
site: "datadoghq.com" # Or eu.datadoghq.com for EU  
host_metadata:  
enabled: true  
hostname_source: "config_or_system"
```

Keep existing exporters commented

jaeger:

endpoint: jaeger:14250

```
service:  
pipelines:  
traces:  
receivers: [otlp]  
processors: [batch, memory_limiter]  
exporters: [datadog] # Changed from jaeger
```

```
metrics:  
receivers: [otlp, prometheus]  
processors: [batch, memory_limiter]  
exporters: [datadog] # Changed from prometheus
```

Apply Changes:

Set Datadog API key

```
export DD_API_KEY="your-api-key-here"
```

Restart collector

```
sudo docker-compose restart otelcol
```

Verify connection

```
sudo docker logs otelcol | grep "datadog"
```

Resource Decorators

Add metadata to all telemetry data for better organization:

```
processors:  
resource:  
attributes:  
- key: environment  
value: production  
action: upsert  
- key: region  
value: us-east-1  
action: upsert  
- key: team  
value: platform-engineering  
action: insert  
- key: version  
value: "1.0.0"  
action: upsert  
  
service:  
pipelines:  
traces:  
receivers: [otlp]  
processors: [resource, batch]  
exporters: [jaeger]
```

Sampling Strategies

Reduce data volume while retaining important traces[5]:

Head-Based Sampling

Sample at trace start (simple but less intelligent):

```
processors:  
probabilistic_sampler:  
sampling_percentage: 10 # Keep 10% of traces  
  
service:  
pipelines:  
traces:  
receivers: [otlp]
```

```
processors: [probabilistic_sampler, batch]
exporters: [jaeger]
```

Tail-Based Sampling

Sample after trace completes (requires buffering):

```
processors:
tail_sampling:
policies:
# Always sample error traces
- name: error_traces
type: status_code
status_code:
status_codes: [ERROR]
```

```
# Sample traces with high latency
```

```
- name: high_latency
type: latency
latency:
threshold_ms: 5000
```

```
# Probabilistic fallback
```

```
- name: probabilisticFallback
type: probabilistic
probabilistic:
sampling_percentage: 5
```

```
service:
pipelines:
traces:
receivers: [otlp]
processors: [tail_sampling, batch]
exporters: [jaeger]
```

Memory Management

Prevent resource exhaustion in the collector:

```
processors:
memory_limiter:
check_interval: 1s # Check every 1 second
limit_mib: 512 # Max 512MB of RAM
spike_limit_mib: 128 # Allow 128MB spike before dropping

batch:
send_batch_size: 1024 # Send after collecting 1024 spans
```

```
timeout: 5s # Or after 5 seconds
send_batch_max_size: 2048

service:
pipelines:
traces:
receivers: [otlp]
processors: [memory_limiter, batch]
exporters: [jaeger]
```

Troubleshooting Guide

Common Issues and Solutions

1. Kafka Service Failing

Symptom: Kafka container exits immediately

Cause: Insufficient RAM

Solutions:

Check available RAM

```
free -h
```

Reduce services using docker-compose override

```
cat > docker-compose.override.yml << EOF
version: '3'
services:
kafka:
environment:
KAFKA_HEAP_OPTS: "-Xms256M -Xmx256M" # Reduce memory
EOF
```

Restart services

```
sudo docker-compose restart kafka
```

2. Port Already in Use

Symptom: Error: bind: address already in use

Cause: Port 80, 8080, 4317, etc. already in use

Solution:

Find process using port 8080

```
sudo lsof -i :8080
```

Kill the process

```
sudo kill -9 <PID>
```

Or use different port in docker-compose

```
docker-compose up -d -p 8081:8080
```

3. Services Not Communicating

Symptom: "Connection refused" errors between services

Cause: Services not on same Docker network or misconfigured endpoints

Solutions:

Check network

```
sudo docker network ls | grep opentelemetry
```

Inspect network

```
sudo docker network inspect opentelemetry-demo
```

Verify all services on same network

```
sudo docker inspect <service_name> | grep NetworkSettings
```

Check service endpoint configuration

```
sudo docker exec <service_name> env | grep OTEL
```

4. No Data in Grafana

Symptom: Dashboards show "no data"

Cause: Collector not receiving or exporting data

Solutions:

Check collector logs

```
sudo docker logs otelcol | tail -50
```

Check for errors

```
sudo docker logs otelcol | grep -i error
```

Verify data flow

```
sudo docker exec otelcol curl -s http://localhost:8889/metrics | head -20
```

Check receiver connectivity

```
sudo docker logs frontend | grep otelcol
```

5. High Memory/CPU Usage

Symptom: Instance becomes slow or unresponsive

Causes: Excessive data collection, sampling not configured

Solutions:

Monitor resource usage

```
watch -n 1 'free -h; echo "---"; top -bn1 | head -15'
```

Enable sampling

```
nano src/otelcollector/otel-collector-config.yml
```

Add probabilistic sampler (see Advanced Configuration)

Restart services

```
sudo docker-compose restart otelcol
```

Debugging Commands

View all logs

```
sudo docker-compose logs -f
```

View specific service logs

```
sudo docker logs -f <service_name>
```

Get into container

```
sudo docker exec -it <service_name> /bin/bash
```

Check metrics endpoint

```
curl http://localhost:8889/metrics
```

Check Prometheus scrape status

```
curl http://localhost:9090/api/v1/targets
```

Verify Jaeger is receiving traces

```
curl http://localhost:16686/api/services
```

Check health of all services

```
sudo docker-compose ps
```

Cleanup and Restart

Stop all services

```
sudo docker-compose down
```

Remove volumes and data

```
sudo docker-compose down -v
```

Clean up images (careful, removes all local images)

```
sudo docker image prune -a
```

Full reset

```
sudo docker system prune -a --volumes
```

Restart fresh

```
sudo docker-compose up -d
```

Best Practices

Implementation Best Practices

1. Semantic Conventions

Follow OpenTelemetry semantic conventions for consistency[6]:

Service Names:

```
service.name: "payment-service"
service.namespace: "production"
service.version: "1.0.0"
```

HTTP Attributes:

```
http.method: "POST"
http.url: "https://api.example.com/v1/payments"
http.status_code: 200
http.response_content_length: 1024
```

Database Operations:

```
db.system: "postgresql"
db.name: "orders_db"
db.operation: "SELECT"
db.statement: "SELECT * FROM orders WHERE id = ?"
```

2. Attribute Management

Minimize Cardinality[7]:

✓ GOOD:

```
environment: "production" (low cardinality)
http.method: "GET" (fixed values)
service.name: "api-service" (bounded)
```

✗ PROBLEMATIC:

```
userid: "12345678" (high cardinality → millions of values)
```

```
request.id: "uuid..." (unbounded cardinality)
customer.email: "user@example.com" (unbounded)
```

Resource Attributes:

```
processors:
resource:
attributes:
- key: environment
  value: production
- key: deployment
  value: kubernetes
- key: version
  value: "1.2.3"
```

3. Sensitive Data Handling

Never collect:

- Personal Identifiable Information (PII)
- Credit card numbers
- API keys or secrets
- Passwords

Redaction Strategy:

```
processors:
attributes:
actions:
# Drop sensitive attributes
- key: password
  action: delete
- key: credit_card
  action: delete
- key: api_key
  action: delete
```

```
# Hash user IDs for cardinality
- key: user.id
  action: hash_sha256
```

4. Instrumentation Strategy

Prioritize Critical Components:

1. **API Endpoints** (entry points to system)
2. **Database Operations** (performance bottlenecks)
3. **Message Queues** (async processing)
4. **External Service Calls** (dependencies)
5. **Business Operations** (revenue impact)

Avoid Excessive Detail:

- Don't instrument every function call
- Focus on cross-service boundaries
- Use sampling for verbose operations

5. Alert Configuration

Create Alerts Based on SLOs:

Grafana Alert: Error Rate > 5%

```
expr: |  
(rate(http_requests_total{status="5xx"}[5m]) /  
rate(http_requests_total[5m])) > 0.05  
notification: on-call-team
```

Grafana Alert: P99 Latency > 1 second

```
expr: |  
histogram_quantile(0.99, rate(http_duration_seconds_bucket[5m])) > 1  
notification: dev-team
```

Operational Best Practices

1. Centralized Configuration

Use ConfigMaps in Kubernetes

```
kubectl create configmap otel-config --from-file=config.yaml
```

Or environment variables for Docker

```
docker run -e OTEL_CONFIG_YAML="$(cat config.yaml)" ...
```

2. Health Checks

```
extensions:  
health_check:  
endpoint: ":13133"
```

Kubernetes liveness probe

```
livenessProbe:  
httpGet:  
path: /healthz  
port: 13133  
initialDelaySeconds: 5  
periodSeconds: 10
```

3. Rate Limiting and Backpressure

```
processors:  
batch:  
  send_batch_size: 1024  
  timeout: 5s  
  
exporters:  
otlp:  
  endpoint: backend:4317  
  sending_queue:  
    queue_size: 5000 # Buffer up to 5000 items  
    storage: memory # Use memory storage  
  retry_on_failure:  
    enabled: true  
    initial_interval: 5s  
    max_interval: 30s
```

4. Testing and Validation

Test Observability:

Generate test data

```
docker run -it grafana/loki-canary:latest  
--url=http://localhost:3100
```

Test metrics collection

```
curl -X POST http://localhost:4318/v1/metrics  
-d @test_metrics.json
```

Validate trace collection

```
curl -X POST http://localhost:4318/v1/traces  
-d @test_traces.json
```

5. Documentation and Runbooks

Create Incident Runbooks:

Scenario	Runbook
High Error Rate	Check error logs → Review traces → Check deployments
High Latency	Check CPU/Memory → Review database queries → Check network
Service Unavailable	Check logs → Verify network connectivity → Restart service
Data Gaps	Check collector health → Verify receiver configuration

Conclusion

This OpenTelemetry observability project demonstrates enterprise-grade monitoring capabilities using industry-standard tools and practices. By completing this guide, you have:

Achievements

- ✓ Deployed a production-grade microservice application
- ✓ Understood observability's three pillars: metrics, logs, and traces
- ✓ Configured OpenTelemetry Collector for telemetry collection
- ✓ Integrated multiple observability backends (Prometheus, Grafana, Jaeger)
- ✓ Performed root cause analysis using distributed traces
- ✓ Implemented error injection and scenario testing
- ✓ Mastered advanced configuration and troubleshooting

Career Impact

This project significantly enhances your DevOps and Site Reliability Engineering (SRE) credentials:

- **Portfolio Piece:** Add to GitHub and resume
- **Interview Preparation:** Demonstrates hands-on observability expertise
- **Promotion Candidate:** Shows advanced infrastructure knowledge
- **Team Value:** Become observability expert for your organization

Next Steps

1. **Document Your Learning:** Write technical blog post on your experience
2. **Share on LinkedIn:** Tag @Cloud Champ and @Open Telemetry
3. **Customize the Demo:** Integrate with your preferred observability backend
4. **Scale to Production:** Adapt patterns to your organization's environment
5. **Contribute:** Submit improvements to the open-source project

References

- [1] OpenTelemetry Project. (2024). OpenTelemetry Documentation. Retrieved from <https://opentelemetry.io/>
- [2] Cloud Champ. (2024, August 22). Observability DevOps Project - OpenTelemetry Demo Application [Video]. YouTube.
- [3] CrowdStrike. (2024, October). The Three Pillars of Observability: Logs, Metrics, and Traces. Retrieved from <https://www.crowdstrike.com/>
- [4] OpenTelemetry. (2024). Collector Configuration. Retrieved from <https://opentelemetry.io/docs/collector/configuration/>
- [5] Better Stack. (2025, March 2). Essential OpenTelemetry Best Practices for Robust Observability. Retrieved from <https://betterstack.com/>
- [6] OpenTelemetry. (2024). Semantic Conventions. Retrieved from <https://opentelemetry.io/docs/specs/semconv/>
- [7] Grafana. (2024, November). OpenTelemetry Best Practices. Retrieved from <https://grafana.com/blog/2023/12/18/>

Document Version: 1.0

Last Updated: December 4, 2024

Author: DevOps Engineering Team

Status: Complete and Ready for Implementation